


Apache Spark

COMP9313: Big Data Management

Motivation of Spark

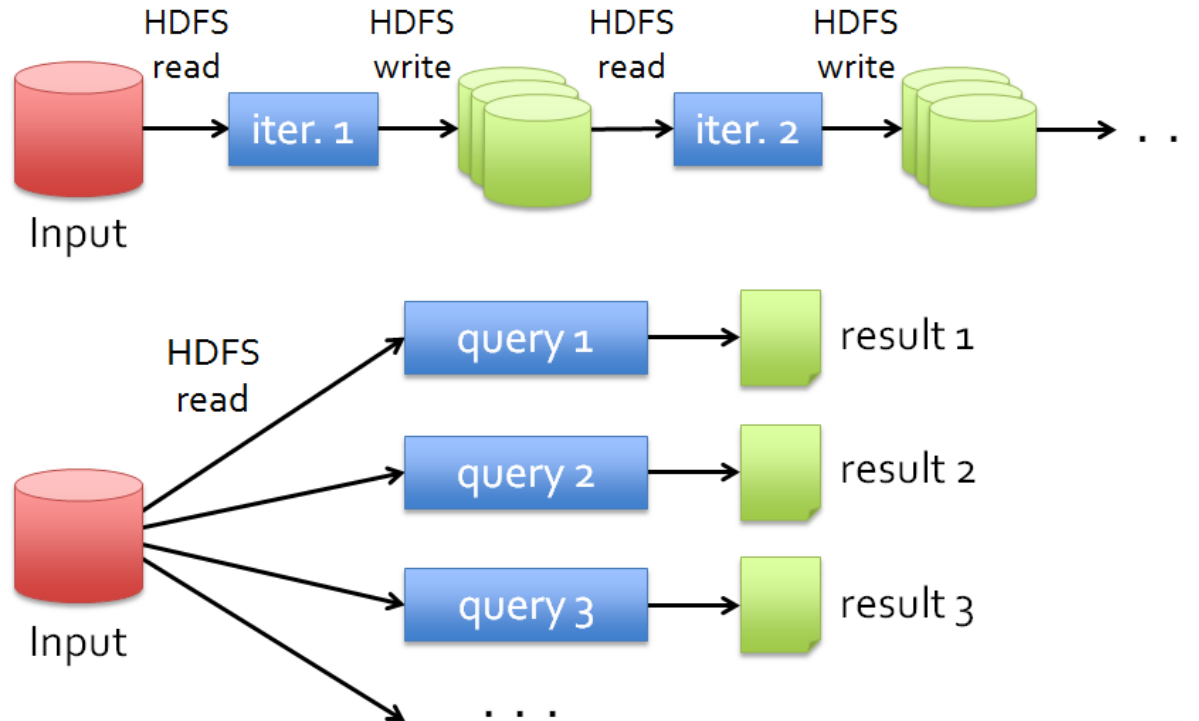
- MapReduce greatly simplified big data analysis on large, unreliable clusters. It is great at one-pass computation
- But as soon as it got popular, users wanted more:
 - More **complex**, multi-pass analytics (e.g. ML, graphs)
 - More **interactive** ad-hoc queries
 - More **real-time** stream processing
- All 3 need faster **data sharing** across parallel jobs
 - One reaction: specialized models for some of these apps, e.g.,
 - Pregel (graph processing)
 - Storm (stream processing)

Limitations of MapReduce



- As a general programming model:
 - It is more suitable for one-pass computation on a large dataset
 - Hard to compose and nest multiple operations
 - No means of expressing iterative operations
- As implemented in Hadoop
 - All datasets are read from disk, then stored back on to disk
 - All data is (usually) triple-replicated for reliability

Data Sharing in MapReduce



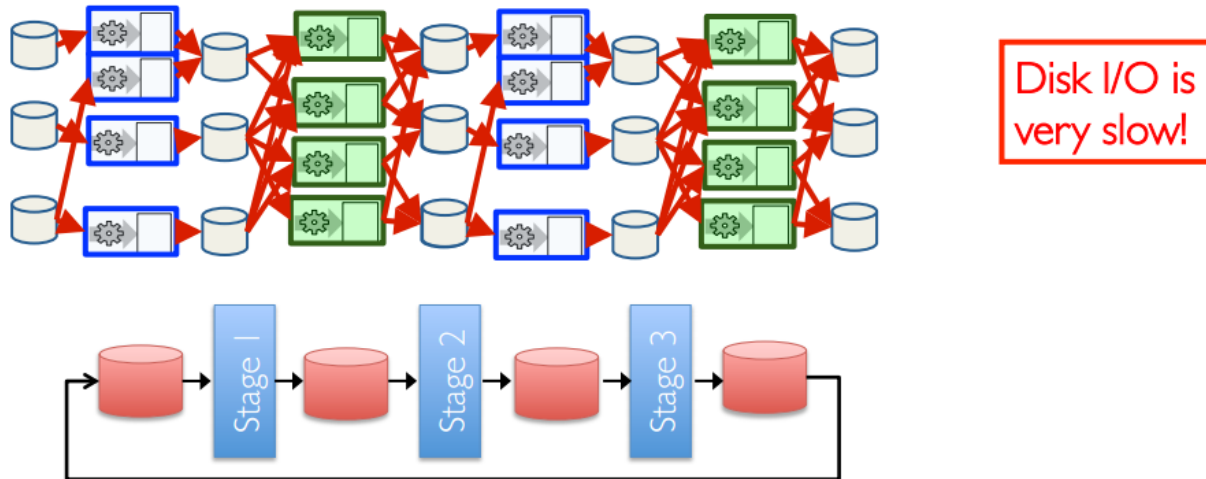
Slow due to replication, serialization, and disk IO

- Complex apps, streaming, and interactive queries all need one thing that MapReduce lacks:

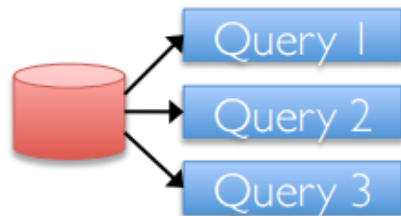
Efficient primitives for **data sharing**

Data Sharing in MapReduce

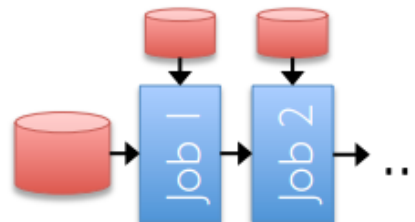
- Iterative jobs involve a lot of disk I/O for each repetition



- Interactive queries and online processing involves lots of disk I/O



Interactive mining




Stream processing

Hardware for Big Data



Lots of hard drives



Lots of CPUs




And lots of memory!

Goals of Spark

- Keep more data in-memory to improve the performance!
- Extend the MapReduce model to better support two common classes of analytics apps:
 - Iterative algorithms (machine learning, graphs)
 - Interactive data mining
- Enhance programmability:
 - Introduce rich API libraries
 - More to be done with less Lines of Code


Data Sharing in Spark with RDDs (Resilient Distributed Dataset)





10-100 × faster than network and disk

What is Spark?

- One popular answer to “What’s beyond MapReduce?”
- Open-source engine for large-scale data processing
 - Supports generalized dataflows
 - Written in Scala, with bindings in Java, Python and R
- Brief history:
 - Developed at UC Berkeley AMPLab in 2009
 - Open-sourced in 2010
 - Became top-level Apache project in February 2014
 - Commercial support provided by DataBricks

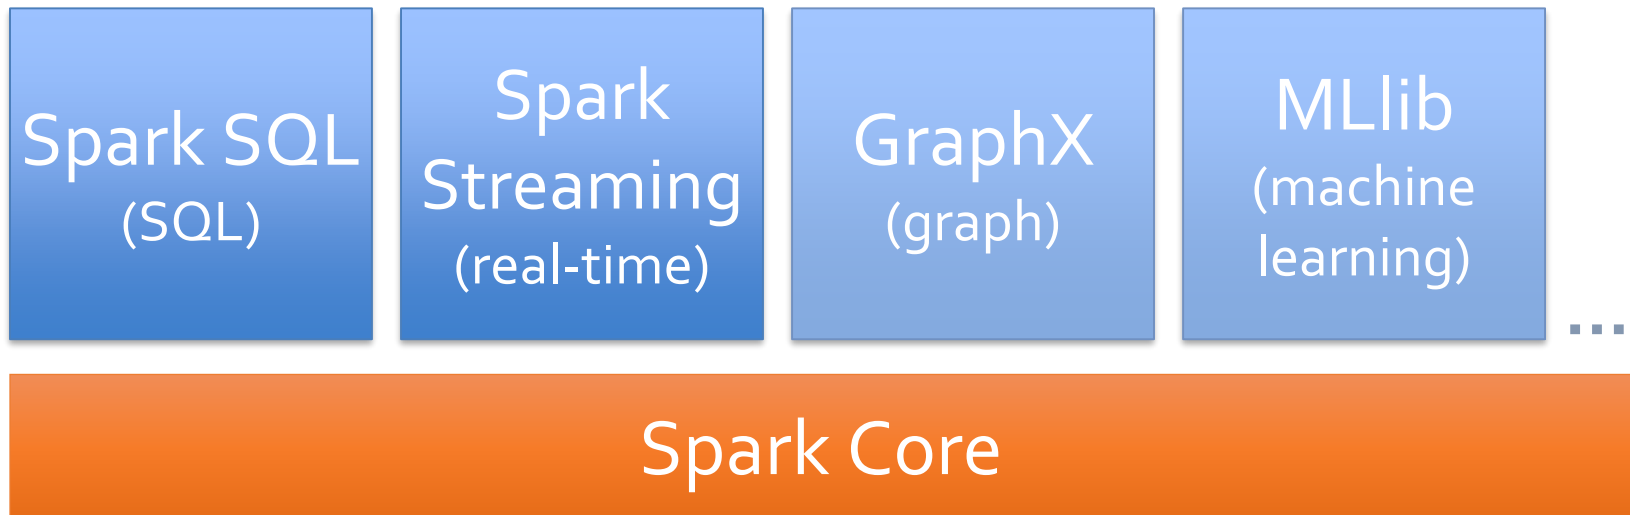


What is Spark?

- Fast and expressive cluster computing system interoperable with Apache Hadoop
- Improves efficiency through:
 - **In-memory** computing primitives
 - General computation graphs Up to 100 × faster
(10 × on disk)
- Improves usability through:
 - Rich APIs in Scala, Java, Python, R
 - Interactive shell Often 5 × less code
- **Spark is not**
 - a modified version of Hadoop
 - dependent on Hadoop because it has its own cluster management (Spark can use Hadoop YARN and HDFS)

What is Spark?

- Spark is the basis of a wide set of projects in the Berkeley Data Analytics Stack (BDAS)



- Spark SQL (SQL on Spark)
- Spark Streaming (stream processing)
- GraphX (graph processing)
- MLlib (machine learning library)


Spark Running Modes

You can run Spark using:

- Its [standalone cluster mode](#)
- On [Hadoop YARN](#)
- On Apache [Mesos](#)
- On [Kubernetes](#)
- Or on the Cloud (e.g., DataBricks).

Data Sources

- Local Files
 - `file:///opt/httpd/logs/access_log`
- Amazon S3
- Hadoop Distributed Filesystem
 - Regular files, sequence files, any other Hadoop InputFormat
- HBase, Cassandra, etc.




Spark Ideas


- Expressive computing system, not limited to map-reduce model
- Facilitate system memory
 - avoid saving intermediate results to disk
 - cache data for repetitive queries (e.g. for machine learning)
- Layer an in-memory system on top of Hadoop
- Achieve fault-tolerance by re-execution instead of replication

Spark Cluster

- To use Spark, developers write a driver program that connects to a cluster of worker



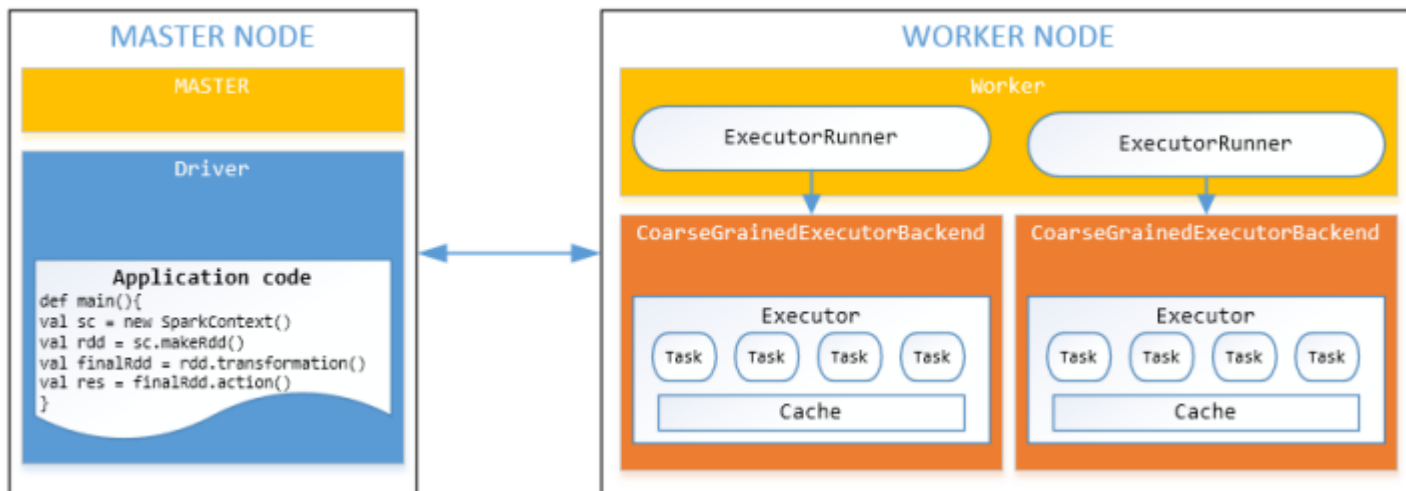
Spark Workflow



- A Spark program first creates a SparkContext object
 - Tells Spark how and where to access a cluster
 - Connects to several types of cluster managers (e.g. YARN, Mesos, or its own manager)
- Cluster manager:
 - Allocates resources across applications
- Spark executor:
 - Runs computations
 - Accesses data storage

Workers Nodes and Executors

- Worker nodes are machines that run executors
 - Host one or multiple Workers
 - One JVM (1 process) per Worker
 - Each Worker can spawn one or more Executors
- Executors run tasks
 - Run in child JVM (1 process)
 - Execute one or more task using threads in a ThreadPool



Introduction to RDDs

Challenge

- Existing Systems

- Existing in-memory storage systems have interfaces based on fine-grained updates
 - Reads and writes to cells in a table
 - E.g., databases, key-value stores, distributed memory
- Requires replicating data or logs across nodes for fault tolerance
 - > expensive!
 - 10-100x slower than memory write

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?

Solution: Resilient Distributed Datasets

- *Resilient Distributed Datasets (RDDs)*
 - Distributed collections of objects that can be cached in memory across cluster
 - Manipulated through parallel operators
 - Automatically recomputed on failure based on lineage
- RDDs can express many parallel algorithms, and capture many current programming models
 - Data flow models: MapReduce, SQL, ...
 - Specialized models for iterative apps: Pregel, ...

What is RDD?

- Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. Matei Zaharia, et al. NSDI'12 ([paper](#))
 - RDD is a **distributed** memory abstraction that lets programmers perform **in-memory** computations on large clusters in a **fault-tolerant** manner
- **Resilient**
 - Fault-tolerant, is able to recompute missing or damaged partitions due to node failures
- **Distributed**
 - Data residing on multiple nodes in a cluster
- **Dataset**
 - A collection of partitioned elements, e.g. tuples or other objects (that represent records of the data you work with)
- RDD is the primary data abstraction in Apache Spark and the core of Spark. It enables operations on collection of elements in parallel

RDD Traits

- **In-Memory**, i.e. data inside RDD is stored in memory as much (size) and long (time) as possible
- **Immutable** or **Read-Only**, i.e. it does not change once created and can only be transformed using transformations to new RDDs
- **Lazy evaluated**, i.e. the data inside RDD is not available or transformed until an action is executed that triggers the execution
- **Cacheable**, i.e. you can hold all the data in a persistent "storage" like memory (default and the most preferred) or disk (the least preferred due to access speed)
- **Parallel**, i.e. process data in parallel
- **Typed**, i.e. values in a RDD have types, e.g. `RDD[Long]` or `RDD[(Int, String)]`
- **Partitioned**, i.e. the data inside an RDD is partitioned (split into partitions) and then distributed across nodes in a cluster (one partition per JVM that may or may not correspond to a single node)

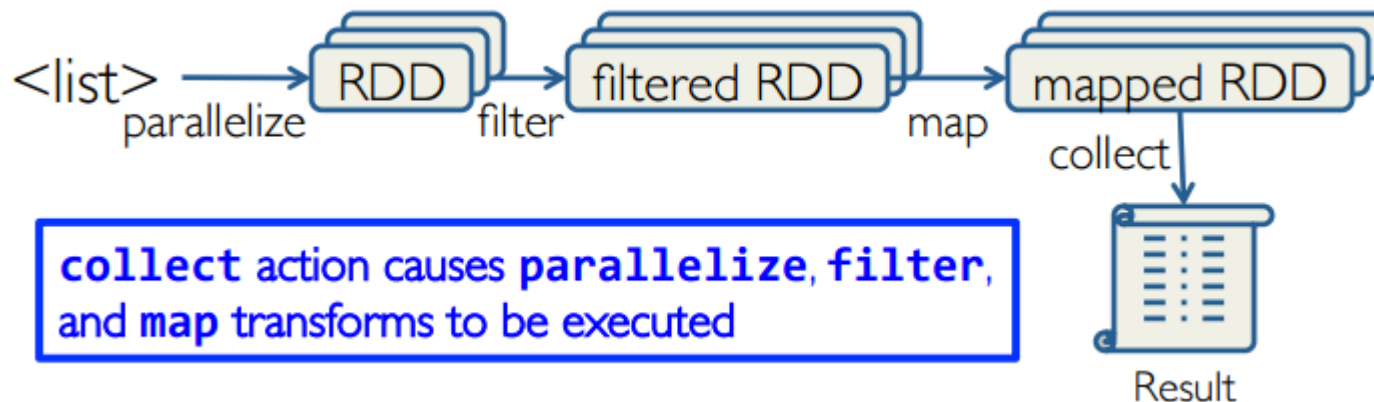
RDD Operations



- **Transformation:** returns a new RDD
 - Nothing gets evaluated when you call a Transformation function, it just takes an RDD and return a new RDD
 - Transformation functions include *map*, *filter*, *flatMap*, *groupByKey*, *reduceByKey*, *aggregateByKey*, *filter*, *join*, etc.
- **Action:** evaluates and returns a new value
 - When an Action function is called on a RDD object, all the data processing queries are computed at that time and the result value is returned
 - Action operations include *reduce*, *collect*, *count*, *first*, *take*, *countByKey*, *foreach*, *saveAsTextFile*, etc.

Working with RDDs

- Create an RDD from a data source
 - by parallelizing existing collections (lists or arrays)
 - by transforming an existing RDDs
 - from files in HDFS or any other storage system
- Apply transformations to an RDD: e.g., map, filter
- Apply actions to an RDD: e.g., collect, count



- Users can control two other aspects:
 - Persistence
 - Partitioning

Creating RDDs

- From HDFS, text files, Amazon S3, Apache HBase, SequenceFiles, any other Hadoop InputFormat
- Creating an RDD from a File
 - RDD distributed in 4 partitions
 - Elements are lines of input

```
JavaRDD<String> distFile = sc.textFile("data.txt",4);
```

- Turn a collection into an RDD

```
List<Integer> data = Arrays.asList(1, 2, 3, 4, 5);  
JavaRDD<Integer> distData = sc.parallelize(data);
```

Spark Transformations

- Create new datasets from an existing one
- Use lazy evaluation: Results not computed right away – instead Spark remembers set of transformations applied to base dataset
 - Spark optimizes the required calculations
 - Spark recovers from failures
- Some transformation functions

Transformation	Meaning
map (<i>func</i>)	Return a new distributed dataset formed by passing each element of the source through a function <i>func</i> .
filter (<i>func</i>)	Return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true.
flatMap (<i>func</i>)	Similar to map, but each input item can be mapped to 0 or more output items (so <i>func</i> should return a Seq rather than a single item).

Spark Actions

- Cause Spark to execute recipe to transform source
- Mechanism for getting results out of Spark
- Some action functions

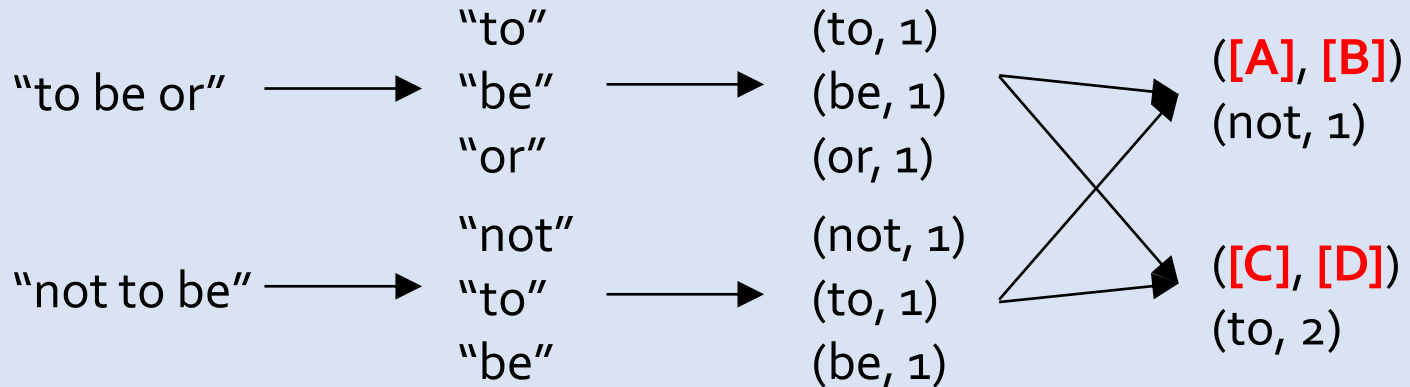
Action	Meaning
<code>reduce(func)</code>	Aggregate the elements of the dataset using a function <i>func</i> (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.
<code>collect()</code>	Return all the elements of the dataset as an array at the driver program. This is usually useful after a filter or other operation that returns a sufficiently small subset of the data.
<code>count()</code>	Return the number of elements in the dataset.

- Example: `counts.saveAsTextFile("hdfs://...");`

Word Count in Spark

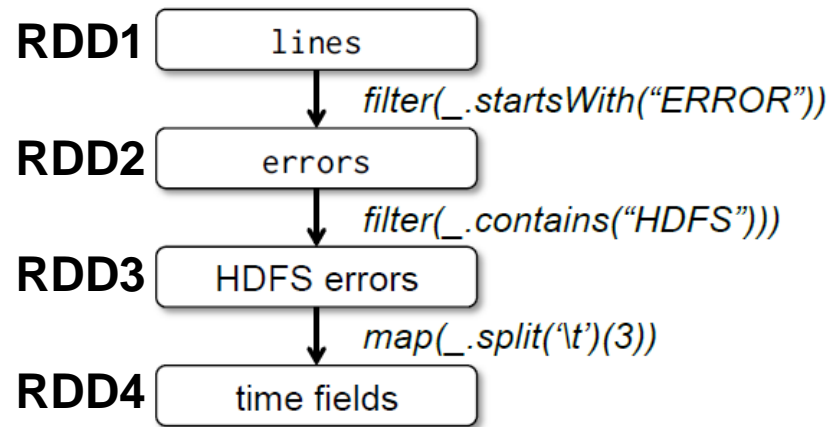
In this example we can use a few transformations to build a dataset of (String, Int) pairs called counts and then save it to a file

```
JavaRDD<String> textFile = sc.textFile("hdfs://...");
JavaPairRDD<String, Integer> counts = textFile
    .flatMap(s -> Arrays.asList(s.split(" ")).iterator())
    .mapToPair(word -> new Tuple2<>(word, 1))
    .reduceByKey((a, b) -> a + b);
counts.saveAsTextFile("hdfs://...");
```



Lineage Graph

- RDD is lazy in nature. It means a series of transformations are performed on an RDD, which is not even evaluated immediately
- RDDs keep track of lineage
- RDD has enough information about how it was derived from to compute its partitions from data in stable storage



- Example:
 - If a partition of errors is lost, Spark rebuilds it by applying a filter on only the corresponding partition of lines
 - Partitions can be recomputed in parallel on different nodes, without having to roll back the whole program

SparkContext

- SparkContext is the entry point to Spark for a Spark application
- Once a SparkContext instance is created you can use it to
 - Create RDDs
 - Create accumulators
 - Create broadcast variables
 - Access Spark services and run jobs
- A Spark context is essentially a client of Spark's execution environment and acts as the *master of your Spark application*
- The first thing a Spark program must do is to create a SparkContext object, which tells Spark how to access a cluster
- In the Spark shell, a special interpreter-aware SparkContext is already created for you

RDD Persistence: Cache/Persist

- One of the most important capabilities in Spark is *persisting* (or *caching*) a dataset in memory across operations
- When you persist an RDD, each node stores any partitions of it. You can reuse it in other actions on that dataset
- Each persisted RDD can be stored using a different *storage level*, e.g.
 - MEMORY_ONLY:
 - Store RDD as deserialized Java objects in the JVM
 - If the RDD does not fit in memory, some partitions will not be cached and will be recomputed when they're needed
 - This is the default level
 - MEMORY_AND_DISK:
 - If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed
- `cache()` = `persist(StorageLevel.MEMORY_ONLY)`

Why Persisting RDD?

- Persist will tell Spark to cache the data in memory, to reduce the data loading cost for further actions on the same data
- `errors.persist()` will do nothing. It is a lazy operation. But now the RDD says "read this file and then cache the contents". The action will trigger computation and data caching

Spark Key-Value RDDs

- Similar to Map Reduce, Spark supports Key-Value pairs
- Each element of a *Pair RDD* is a pair tuple
- Some Key-Value transformation functions:

groupByKey([*numPartitions*])

When called on a dataset of (K, V) pairs, returns a dataset of (K, Iterable<V>) pairs.

Note: If you are grouping in order to perform an aggregation (such as a sum or average) over each key, using `reduceByKey` or `aggregateByKey` will yield much better performance.

Note: By default, the level of parallelism in the output depends on the number of partitions of the parent RDD. You can pass an optional `numPartitions` argument to set a different number of tasks.

reduceByKey(*func*, [*numPartitions*])

When called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs where the values for each key are aggregated using the given reduce function *func*, which must be of type (V,V) => V. Like in `groupByKey`, the number of reduce tasks is configurable through an optional second argument.

Questions?